

# Edward Vaisman

✉ [vaismanedward@gmail.com](mailto:vaismanedward@gmail.com) | 💻 [eddyv](#) | 🌐 [edwardvaisman.ca](http://edwardvaisman.ca) | 📺 [edwardvaisman](#) | ☎ (+1) 647-219-6236

## Professional Summary

Software engineer with 8+ years of expertise specializing in distributed systems and event-driven architectures. Proven track record in designing and implementing enterprise-scale Kafka ecosystems, cloud-native solutions, and real-time data pipelines across various cloud providers. Known for architecting sophisticated systems - from protocol-level engineering to AI-powered applications - while championing technical excellence and automating workflows.

## Professional Experience



### Confluent

Staff Partner Innovation Engineer

Toronto, ON

Oct 2022 - Current

- Designed & built a Kafka proxy to interpret Kafka's wire protocol, bridging clients and brokers to unlock advanced Confluent Cloud use cases via pluggable modules (Authentication Transformers, Interceptors).
- Developed a DynamoDB Source Connector for a client, later promoted to a fully managed source connector.
- Enhanced consumption tracking methodologies, driving insights into confluent cloud usage and billing.
- Delivered an end-to-end Generative AI and Retrieval-Augmented Generation (RAG) application, integrating Confluent, Apache Flink, AI Model Inference, MongoDB Atlas, and a suite of AWS services (including Lambda, API Gateway, Secrets Manager, and S3), showcased at AWS re:Invent and webinars as a go-to-market initiative, strengthening strategic partnerships through co-marketing and technical collaboration.
- Engineering cutting-edge Model Context Protocol (MCP) architectures to enhance LLM capabilities and RAG system integration.

Kafka Flink KStream Model Context Protocol Terraform Java Vert.x GenAI AWS GCP Azure



### Wavelo (Tucows)

Senior Staff Engineer (Kafka)

Toronto, ON

Nov 2020 - Oct 2022

- Architected and secured a robust, multi-datacenter Kafka infrastructure, including add-ons like Confluent Schema Registry, MirrorMaker2, and Kafka Connect.
- Developed automation for provisioning, maintaining, and monitoring Kafka environments using the HashiCorp stack.
- Designed, implemented, and operated multi-datacenter Kafka infrastructure along-side kafka add-ons such as Confluent Schema Registry, MirrorMaker2, Kafka Connect.
- Implemented observability solutions, creating real-time monitoring dashboards to track Kafka cluster connectivity, security, and performance. (DataDog, Elasticsearch, Prometheus, Grafana)
- Built self-serve GitOps tools to increase speed and productivity of developers using the kafka platform. (Topic Management, ACL Management, Kafka Connect Cluster Deployment, KSQL Deployment, Self-Hosted GitHub Runners, etc...)
- Increased our PaaS offerings by building an in-house CLI to interact with our private cloud infrastructure to deal with secret management, volume management, application deployment. (Typer CLI, Nomad APIs, Vault APIs)
- Optimized Kafka performance by identifying bottlenecks and implementing throughput enhancements.
- Ensured system reliability through rigorous disaster recovery testing and routine backup and restoration exercises.
- Educated staff on Kafka architecture, Schema Registry, and event-driven design principles, fostering company-wide adoption of best practices.

Kafka Terraform Nomad Vault SaltStack Docker GitHub Actions Prometheus Grafana DataDog Elasticsearch



### Bank Of Montreal (BMO)

Senior Application Developer

Toronto, ON

Sept 2018 - Oct 2020

- Played a key role in advancing the bank's Customer Intelligence Platform by delivering a near real-time solution to provide the business with a unified view of customer interactions and enhance customer experience.
- Onboarded, trained, and mentored developers, ensuring comfortability and expertise in distributed systems, Git Flow, Dockerized deployments, and event-driven architecture.
- Designed and implemented a secure, field-level encrypted real-time data pipeline enabling seamless data transfer between on-premises and cloud-based Kafka clusters (AWS).
- Developed and optimized real-time streaming data pipelines to facilitate reliable data exchange across systems and applications.
- Created real-time streaming applications to transform and react to data streams, driving actionable insights.
- Automated workflows and integrations to enhance operational efficiency and effectiveness.
- Produced high-quality software implementations, adhering to stringent standards for security, logging, error handling, performance, and other non-functional requirements.

Kafka Kafka Streams Java Spring Boot AWS



- Managed web pages, email hosting, and web hosting services to ensure seamless client operations and user satisfaction.
- Designed and developed scalable, responsive, and robust web pages optimized for high performance and conversion rates.
- Collaborated with clients and project managers to define project scope, goals, and deliverables, including SEO optimization strategies.
- Gathered and analyzed user data to enhance website performance and presented actionable insights to clients.
- Delivered detailed software solutions encompassing animations (CSS3, JQuery), scripting (PHP, JavaScript, JQuery, Google Apps Scripts), and CMS development (WordPress).
- Built proficiency in time management by consistently meeting tight deadlines across multiple projects.

HTML CSS JavaScript PHP WordPress JQuery Google Apps Scripts SEO

## Skills

---

<b>Programming</b>	Java, Python, Bash, C#, XNA Game Studio
<b>Web</b>	HTML5, CSS, JavaScript, Tailwindcss, Astro, WordPress, React
<b>Big Data</b>	Confluent, Kafka, Flink, DataHub
<b>Cloud</b>	AWS, GCP, Azure, OpenStack, DigitalOcean
<b>DevOps</b>	GitHub, GitHub Actions, SemaphoreCI, SaltStack, Terraform, Nomad, Vault, Docker, GitLab, Ansible, Maven
<b>Languages</b>	English, Russian.

## Education

---

**York University**  
B.Sc, Spec. Hons. Computer Science (Software Development)

*Toronto, ON*  
*Sept 2013 - April 2018*